

Bias in Facial Recognition Technology

By: Maxwell Yearwood and Christine Marra

CSCI 77800 Ethics

Fall 2022

Overview

In this paper we will discuss:

- History of Facial Recognition
- Modern Facial Recognition Technology
- Current Uses of Facial Recognition Technology
- Shortcomings Facial Recognition Technology
- Recommendations to Improve Facial Recognition Technology
- Legislative Measures to Ensure the Ethical and Transparent Use of Facial Recognition

History Facial Recognition

Each person has certain biological features, a.k.a. biometrics, that are unique to them alone. A person's face is one of several unique biometric identifiers; others being fingerprints, voice and vein patterns in the eye and palm. The earliest known use of biometrics to identify a person dates back to the [Qin Dynasty](#) in China 221B.C. by using a clay finger imprint.

The first use of facial recognition, the ability to identify an individual by their face, as a tool dates back to the late 19th century. England introduced a regular system of [prison photography](#) to make prisoners easier to find in the event of a successful escape and to enable record-sharing with other police stations. The first use of facial recognition as a tool in the United States is credited to the Pinkerton National Detective Agency. Before the establishment of a national law enforcement agency (i.e.FBI), the Pinkerton National Detective Agency investigated interstate crimes involving train and bank robberies. In 1895 they published [Criminal Mug Shot & Information Book](#), an accumulation of 300 mugshots of known thieves. At the time, it was a “must have” for all members of The American Bankers Association.

As Alan Turing, is considered the father of computer science, [Woodrow Bledsoe](#) can be considered the father of facial recognition technology (FRT). In the mid 1960's Bledsoe and two colleagues at [Panoramic Research, Inc. \(PRI\)](#), developed a process of comparing a photo to a database of photos to find a match. They referred to their

project as “man-machine” because the process of identifying facial “landmarks,” which is now referred to as measurement and extraction and done by computers, was done by humans.

The next generation of FRT using algorithms was created in 1991 by MIT researchers Matthew Turk and Alex Pentland. Building upon the idea of [Sirovich and Kirby](#), which was based on the idea that all pictures of the human face to be a weighted sum of a few “key pictures” they coined “eigenpictures.” In a study by [Turk and Pentland](#) they laid out a memory-efficient way to compute the eigenpictures. They also proposed an algorithm on how the face recognition system can operate, including how to update the system to include new faces and how to combine it with a video capture system.

SOURCE

Modern Facial Recognition Technology

Fast forward to the 21st century and the use of facial recognition as a tool has expanded in scope and sophistication. Today, the process of identifying someone by their face is exclusively digital. At the highest level, facial recognition technology relies on artificial intelligence (AI) to learn the patterns of a human face. AI systems use machine learning models to learn from a dataset of human faces. These datasets, (i.e. photos) are compiled using various sources including, but not limited to, social media platforms, DMV records and police files.

The process of identifying a person using facial recognition technology (FRT), takes the following steps: detection, alignment, measurement/extraction, recognition and verification.

Detection and Alignment

All FRT programs must first detect or identify the face in an image or video.

Unfortunately, in most photos, the face is not squarely centered. This poses a problem because to a computer, faces that are turned away from the focal point appear completely different. Bledsoe’s FTR fell short because he was unable to account for rotations or tilts of the face in photos. Today, an algorithm aligns or centers the face to make it consistent with the faces in the database. During this [process](#) a variety of generic face “landmarks” such as the bottom of the chin to the top of the nose, the outsides of the eyes, and different places surrounding the eyes and lips are used. The next step is to then train a deep learning (AI) system to locate these spots on any face and turn them towards the center. This makes the face detection process possible.

Measurement and Extraction

Every face has numerous, distinguishable “landmarks”, the different peaks and valleys that make up facial features. [Facelt](#) defines these landmarks as nodal points. Each human face has approximately 80 nodal points. Some of these measured by the software are:

- Distance between the eyes
- Width of the nose
- Depth of the eye sockets
- The shape of the cheekbones
- The length of the jaw line

A deep learning (AI) system then measures the curves of the face and creates a template.

Recognition

Using the unique measurements just generated, AI algorithms then compare them to the measurements of each face in the given database/dataset. The match will be whichever face in the database comes closest to the measurements of the face in question.

Verification

The last step, also performed by deep learning algorithms, is to match the face with other faces in the database. If the face matches then it is said to be verified, and if unmatched it remains unverified. This verification process can sometimes be complex. The facial images created during the measurement and extraction process are 3-D. Problems arise when the images in the dataset are 2-D.

Current Uses of Facial Recognition Technology

Most people are unaware of the extent to which facial recognition technology pervades their everyday lives. Many of us use FTR to safeguard and access our personal devices. Our phones and tablets come with FRT already installed, so whether knowingly or unwittingly, we use FRT everytime we share and store photos. Many “techies” and “early adopters,” will have the ability to speed through lines at airports, hotels and grocery stores.

Facial recognition technology is currently being used in many public spaces where large groups of people gather and security is a concern, especially airports and stadiums. The trade-off for this convenience and security is a lack of privacy. The use of security cameras in public places is a source of photos for a potential dataset.

The most controversial use of FRT is in law enforcement and community policing. Law enforcement agencies ranging from local police departments to the U.S. Customs and Border Dep't. use FRT to identify and apprehend criminals. Although there are many success stories and enhanced public safety resulting from this technology, it is still imperfect and needs to be used in conjunction with other tools. When used exclusively, without proper checks and balances the results can be devastating. [Nijeer Parks](#) and [Robert Williams](#) highlight when law enforcement officers used FRT irresponsibly.

Facial Recognition Technology Shortcomings


Given our highly partisan and divided society, one thing most of us can agree upon is that the major drawback of facial recognition technology is the occurrence of false positive results. The range of false positive error rates in one study was between 3 errors out of 100,000 queries (0.003 percent) in optimal conditions to 3 errors out of every 1,000 queries (0.3 percent).

Factors to Consider:

- Quality of photo
- Performance of the algorithm
- Size of database
- Type of comparison; one-to-one or one-to-many

[Source](#)

The existence of false positive results is referred to as implicit bias because it happens most often to women of color.



[SOURCE](#)

Recommendations to Improve Facial Recognition Technology

There are a few ways to improve the use of facial recognition technology. One action that can be taken to improve the verification results is to improve the datasets used FRT algorithms. Underexposed and poor quality photos should be removed from datasets as it makes the performance of the results of the facial extraction and measurement process less reliable. More varied photos need to be added to the datasets in the AI algorithms. It is noteworthy to mention that China does not experience the same level of inaccuracy in their FR systems as we do here in the U.S. One possible conclusion that can be drawn is that since the Chinese population is more homogeneous than that of the U.S., all of the photos in their datasets will have darker skin tones. Therefore, since the FR algorithms use machine learning, the more datasets, the more accurate the results.

Algorithms

Algorithms and machine learning (AI) form the foundation Facial Recognition Technology. There are currently hundreds of biometric authentication recognition algorithms functional and patented within the USA. Therefore, improved FRT, will need to come from improved algorithms.

Legislative Measures to Ensure the Ethical and Transparent Use of Facial Recognition Technology

[Illinois Biometric Information Privacy Act \(BIPA\)](#).

[NYS Assembly Bill A6787D](#) - banned the use of biometric surveillance technology in schools **until** the New York State Education Department (NYSED) issues a report on the risks and benefits of this technology in schools and the Commissioner of Education authorizes its use.

To Be Developed

Set national standards for algorithms

research literature is often opaque, but public discussion of FRT must be better informed. FRT is improving rapidly, and any critique based on data from even a few years ago runs the risk of being entirely wrong. Determining ground truth on issues of accuracy and bias is essential for good policy.

Sources

<https://bioconnect.com/2021/12/08/a-brief-history-of-biometrics/>

<https://www.smithsonianmag.com/history/first-case-where-fingerprints-were-used-evidence-180970883/>

<https://leg.mt.gov/content/Committees/Interim/2021-2022/Economic%20Affairs/Meetings/September%202021/facial-recognition-technology.pdf>

Legal concerns

<https://www.csis.org/analysis/facial-recognition-technology-responsible-use-principles-and-legislative-landscape>

<https://www.mitre.org/sites/default/files/2021-11/biometric-face-recognition-references-policymakers.pdf>

How does Facial Recognition work?

<https://www.csis.org/analysis/how-does-facial-recognition-work>

<https://www.thalesgroup.com/en/markets/digital-identity-and-security/government/biometrics/facial-recognition>

Facial Recognition is everywhere.

<https://www.nytimes.com/wirecutter/blog/how-facial-recognition-works/>

Safety of Facial Recognition

<https://www.electronicid.eu/en/blog/post/face-recognition/en>

Code Portion of Final Project

How does facial recognition work?

Facial recognition technology takes place in two main ways: facial detection and facial identification. Facial detection is an essential first step in facial recognition. In this step, the software uses a mathematical formula to reduce the many features of your face into tiny bits of information. Some algorithms use cascades to repeatedly scan the features of a face. This step constitutes one's facial signature. Once the face is detected it is then compared to a known database of faces. The unknown face is then matched to a known face. This completes the identification phase and thus facial recognition is complete. These steps are illustrated in our code in Appendix A and B.

Our simulations of facial recognition are illustrated in two ways. The code in appendix A illustrates facial recognition by comparing a known database of faces housed on a hard drive against an unknown face housed in another folder on the same hard drive. The code in appendix B illustrates facial recognition technology in real-time by using a webcam and a phone. An unknown image on the cell phone is held up against the webcam at various distances and angles. The known images are stored in a folder on the computer's hard drive. Once a face is detected a rectangle is drawn around the face along with the name(s) that match the image shown on the phone. The code displays "unknown" if there is no match or an error message if no image is detected. This completes the facial recognition process.

We used images of Keanu Reeves, Will Ferrel, and Idris Elba in this simulation. Our simulations confirm that this technology detects and identifies male caucasoid faces accurately but often fails to do so for some known dark skinned male images. These findings are in accordance with the literature on facial recognition technology.

Appendix A

```
import cv2
import face_recognition

#This algorithm detects all images for Keanu Reeves and Will Ferrel in my database
# but could not detect some of Idris Elba's images. Instead it
# displays the error: " IndexError: list index out of range error"
# Why is this so? According to Stackoverflow
#
(https://stackoverflow.com/questions/59919993/indexerror-list-index-out-of-range-face-recognition)
# this algorithm displays the error: " IndexError: list index out of range error"
# because it could not detect a given image. Hence one ethics issue in facial recognition.
#Some experts claim that facial recognition technology (sometimes) misidentifies dark skinned
people

# Accessing and encoding known image
known_image = cv2.imread("C:/Users/Tashema Bholanath/Documents/images/Idris_Elba.jpg ")
rgb_img = cv2.cvtColor(known_image,cv2.COLOR_BGR2RGB)
img_encoding = face_recognition.face_encodings(rgb_img)[0]

#Accessing and encoding unknown image
unknown_image = cv2.imread("C:/Users/Tashema Bholanath/Documents/images/Idris
Elba_2.jpg ")
rgb_img_2 = cv2.cvtColor(unknown_image,cv2.COLOR_BGR2RGB)
img_encoding_2 = face_recognition.face_encodings(rgb_img_2)[0]

# Comparing images to test if known and unknown images represent the same person.Returns
true if images are the same person.
image_recognition = face_recognition.compare_faces([img_encoding],img_encoding_2)
print("Image Recognized:", image_recognition)

# Displays known and unknown images
cv2.imshow("Img",known_image)
cv2.imshow("Img_2",unknown_image)

# Image is displayed until user presses a key on keyboard
cv2.waitKey(0)
```


Appendix B

This algorithm works by using your webcam and phone to provide real-time facial recognition
to simulate real-time facial recognition used in practice.

#Note Findings: Algorithms normally works fine but either fails to detect or name
an unknown face if the phone being held to the Webcam is either held too far
away or if the phone showing the unknown face is held at an angle. Hence an ethical
issue of misidentification of some unknown faces

```
import cv2
import face_recognition
from simple_facerec_import_SimpleFacerec

# Encode all faces in folder

sfr = SimpleFacerec()
sfr.load_encoding_images("C:/Users/Tashema Bholanath/Documents/images/ ")

# Take Webcam Stream
# Load Camera
cap = cv2.VideoCapture(0)

while True:
    ret, frame = cap.read()

    # Face Location and face detection of known faces
    face_locations, face_names = sfr.detect_known_faces(frame)
    for face_loc, name in zip(face_locations, face_names):
        y1, x1, y2, x2 = face_loc[0], face_loc[1], face_loc[2], face_loc[3]
        #print(face_loc)
    #Show names and rectangle around face identified
    cv2.putText(frame, name, (x1, y1-10), cv2.FONT_HERSHEY_DUPLEX, 1, (0, 0, 0), 2)
    cv2.rectangle(frame, (x1, y1), (x2, y2), (0, 0, 200), 2)

    cv2.imshow("Frame", frame)


    key = cv2.waitKey(1)
    if key == 27:
```

break


cap.release()

cv2.destroyAllWindows()

Appendix C



Keanu_Reeves



Keanu_Reeves_1




Idris_Elba_4




Idris_Elba_3



Idris_Elba_1




Idris_Elba_2



Idris_Elba



Will_Ferrel



Wil_Ferrel_2